

Initial Prioritisation of Unstructured and Semi-Structured Data for AI Use Cases

This document defines the initial priority scope for unstructured and semi-structured data relevant to the current AI use case portfolio. Knowledge-domain prioritisation is being handled separately.

Document intent

Provide a practical first-cut view of the data categories and source platforms that should be prioritised for AI use cases, to support source inventory, capability prioritisation, reference architecture work, and gap assessment.

1. Purpose

This document defines the initial priority scope for **unstructured and semi-structured data** relevant to the current AI use case portfolio.

It is intended to:

- identify the priority data types most relevant to current AI use cases
- identify representative source platforms and systems of record
- support source inventory, capability prioritisation, and reference architecture work
- provide a practical first-cut view for immediate programme use

Note: Knowledge-domain prioritisation is being handled separately and is not the primary focus of this document.

2. Context

The bank is building a trusted, AI-ready data foundation to accelerate decision-making, enable responsible AI solutions, and drive measurable business outcomes.

As part of this, the **Unstructured Data for AI** workstream is focused on establishing a strategic approach for how unstructured and semi-structured data is identified, prioritised, ingested, processed, governed, and made available for AI use cases.

A complete source-level inventory is not yet available across all use cases. Therefore, this first iteration prioritises data by:

- content type / content class
- source platform / system of record where known

This allows the workstream to move forward pragmatically while deeper discovery continues.

3. Prioritisation Approach

Level 1 - Content Type / Content Class

This level identifies what kind of unstructured or semi-structured data is required by the AI use cases. Examples include policy and procedure documents, knowledge articles, reports and guidance documents, contracts and forms, email content and email threads, chat conversations, call recordings, transcripts, scanned documents and images, tickets, cases, work items, and metadata-rich semi-structured records.

Level 2 - Source Platform / System of Record

This level identifies where the data originates, where that is already known. Examples include Confluence, SharePoint, ServiceNow knowledge articles, Jira, email platforms or mailboxes, contact-centre platforms, call-recording systems, content-management systems, document repositories, object or blob storage, and internal application databases exporting semi-structured records.

The two-level approach allows the workstream to distinguish clearly between **what content needs to be supported** and **where that content is sourced from**.

4. Prioritisation Principles

- relevance to active AI use cases
- reuse potential across use cases
- strategic value for AI enablement
- need for standardisation and control
- feasibility for near-term architecture progression

5. Recommended Priority Data Categories

The following data categories are recommended as the initial priority scope for the workstream.

Priority data category	Description	Example content classes	Why prioritised
Documents	Long-form document content used for policy, procedure, knowledge, and guidance.	Policies, procedures, manuals, reports, guidance, contracts	Common across multiple AI use cases and foundational for retrieval and grounding.
Audio and transcripts	Voice-based interactions and derived transcriptions.	Call recordings, voice interactions, transcripts	Important for conversational AI and customer-interaction use cases.
Chat and conversational data	Text-based interactions between users, customers, agents, or bots.	Live chat, chatbot conversations, service conversations	Relevant to customer servicing, insight extraction, and AI assistance.
Email and communications	Message-based communication content.	Emails, threads, attachments	High business value in some use cases, but requires strong governance.
Images and scanned content	Image-based or non-born-digital content.	Scanned documents, screenshots, image artefacts	Important where extraction or OCR is required.
Semi-structured business records	Content with partial structure and business metadata.	Tickets, cases, work items, forms, logs, exported records	Important bridge between structured and unstructured data processing.

6. Representative Source Platform Examples

Where source detail is already known or can be reasonably inferred, the following source-platform types are relevant.

Source platform / system type	Example systems	Typical data types
Document and collaboration platforms	Confluence, SharePoint	Documents, knowledge articles, reports, guidance
Service and case-management platforms	ServiceNow, Jira	Knowledge articles, tickets, cases, work items
Email platforms	Enterprise mail systems	Email content, threads, attachments

Source platform / system type	Example systems	Typical data types
Contact-centre and voice platforms	Call platforms, recording systems	Audio, transcripts, interaction logs
Chat and messaging platforms	Live-chat or bot systems	Chat conversations, transcripts
Content-management platforms	Internal KMS or CMS	Documents, managed knowledge artefacts
Storage platforms	Blob storage, object stores, repositories	Raw files, processed artefacts, exports
Application source systems	Internal business applications	Semi-structured records, metadata-rich extracts

7. Initial Mapping from Current Reference Use Cases

Reference use case	Priority data categories	Example content types	Known / likely source platforms
Codifying the Bank	Documents, semi-structured records	Policies, procedures, controlled documents, supporting artefacts	Document repositories, internal content sources
Customer Journey / Voice AI	Audio and transcripts, chat data, semi-structured records	Call recordings, transcripts, chat interactions, conversation logs	Contact-centre platforms, chat platforms, landing storage
Knowledge Management System	Documents, knowledge articles, managed content	Knowledge artefacts, articles, uploaded documents, metadata-rich content	KMS or CMS platforms, blob/document stores, metadata stores

8. Baseline Observations from Current Reference Use Cases

- multiple unstructured-data patterns already exist across the bank
- current delivery spans document-centric, audio/chat-centric, and knowledge-management-led patterns
- useful capabilities already exist across ingestion, transformation, storage, retrieval, and governance
- current approaches are valuable but fragmented and not yet aligned to a common strategic model

This confirms that the workstream should focus on standardising common patterns, defining shared enterprise capabilities, clarifying source-onboarding and control expectations, and shaping the target reference architecture based on real delivery patterns.

9. Recommended Position

The recommended immediate priority is to proceed with an initial unstructured-data scope based on:

- priority content categories: documents, audio and transcripts, chat and conversational data, email and communications, images and scanned content, and semi-structured business records
- representative source platforms where already known: collaboration and document platforms, service and case-management platforms, email platforms, contact-centre and chat platforms, content-management systems, and storage or repository platforms

This provides a practical basis for source inventory, capability prioritisation, reference architecture scope, gap assessment, and strategic offering identification.

10. Next Steps

- Validate the priority data categories against the current AI use case portfolio.
- Add known source platforms and systems of record where already understood.
- Confirm which source types should be included in the initial source inventory.
- Use the agreed priority scope to support reference architecture development, gap assessment, and strategic-offering definition.